


benchmark training loss rolling std



old benchmark training / fixed exact validation | n=training batches | seed=0

Old benchmark models are separate from main. Teacher-specific loss/std are not shared-ground-truth performance or rankings.